

final_project

wz2631

2023-04-30

Data preparation

```
# draw 2 random samples of 2000 participants
load("./recovery.Rdata")
set.seed(2631)
dat1 <- dat[sample(1:10000, 2000),] %>%
  janitor::clean_names() %>%
  na.omit()
set.seed(2855)
dat2 <- dat[sample(1:10000, 2000),] %>%
  janitor::clean_names() %>%
  na.omit()
dat <- rbind.fill(dat1, dat2) %>%
  dplyr::select(-id) %>%
  unique() %>% mutate( gender=fct_recode(factor(gender),male='1',female='0'),
    race=fct_recode(factor(race),white='1',asian='2',black='3',hispanic='4'),
    smoking=fct_recode(factor(smoking),never='0',former='1',current='2'),
    hypertension=factor(hypertension),
    diabetes=factor(diabetes),
    vaccine=factor(vaccine),
    severity=factor(severity),
    study=factor(study),
    recovery_bin = if_else(recovery_time <= 30, 'f', 's'),
    recovery_bin = factor(recovery_bin)
  )
```

```
#data partition
set.seed(2023)
train_index=createDataPartition(y = dat$recovery_time,
                                p = 0.8,
                                list = FALSE)

train_dat=dat[train_index,]
test_dat=dat[-train_index,]
#training data
train_data=dat[train_index,]
x1 = model.matrix(recovery_time~., data=dat)[train_index,-1]
y1=dat$recovery_time[train_index]
#testing data
test_data=dat[-train_index,]
x2=model.matrix(recovery_time~., data=dat)[-train_index,-1]
y2=dat$recovery_time[-train_index]
```

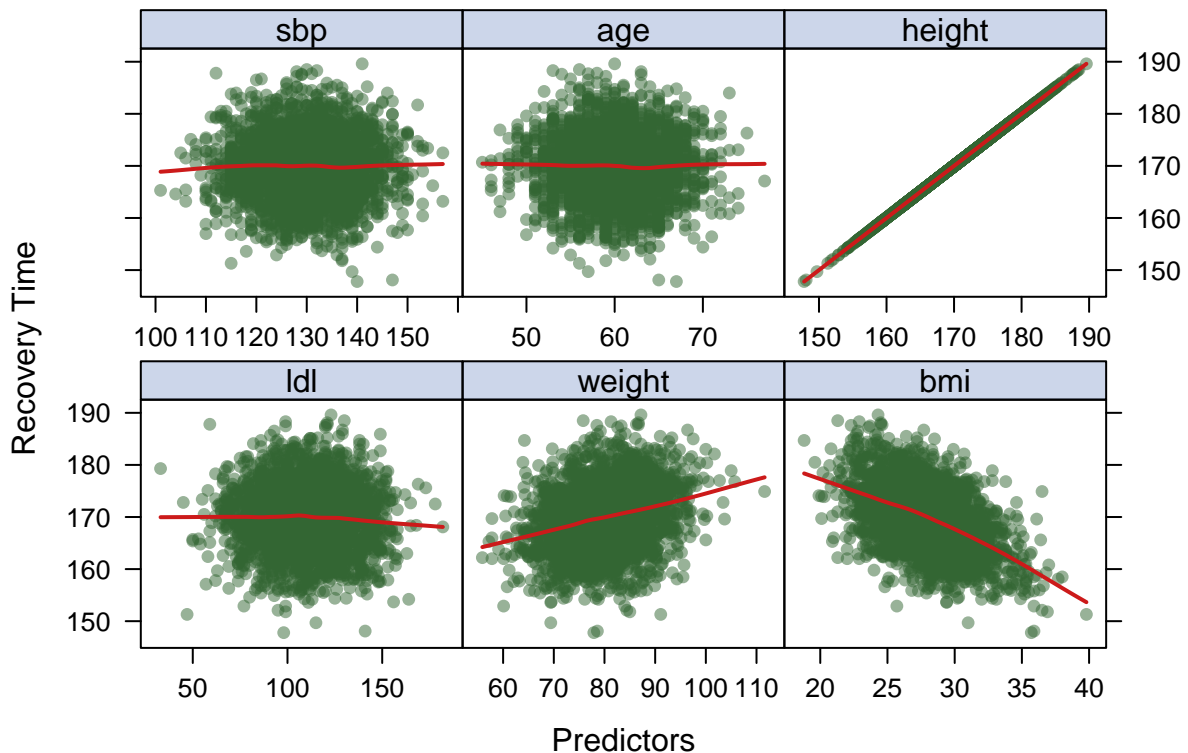
```

#exploratory analysis and data visualization
visualization = train_dat %>%
  mutate(study=case_when(
    study == "A" ~ 1,
    study == "B" ~ 2,
    study == "C" ~ 3
  )) %>%
  dplyr::select(ldl,weight,bmi,sbp,age,height)
non_numeric= sapply(visualization, function(x) !is.numeric(x))
visualization[, non_numeric] = lapply(visualization[, non_numeric], as.numeric)
theme1 = trellis.par.get()
theme1$plot.symbol$col = rgb(.2, .4, .2, .5)
theme1$plot.symbol$pch=16
theme1$plot.line$col=rgb(.8, .1, .1, 1)
theme1$plot.line$lwd=2
theme1$strip.background$col=rgb(.0, .2, .6, .2)
trellis.par.set(theme1)

featurePlot(x = visualization[,1:6],
  y = visualization[,6],
  plot = "scatter",
  span = .5,
  labels = c("Predictors", "Recovery Time"),
  main = "Figure 1. the relationship between predictors and recovery time",
  type = c("p", "smooth"))

```

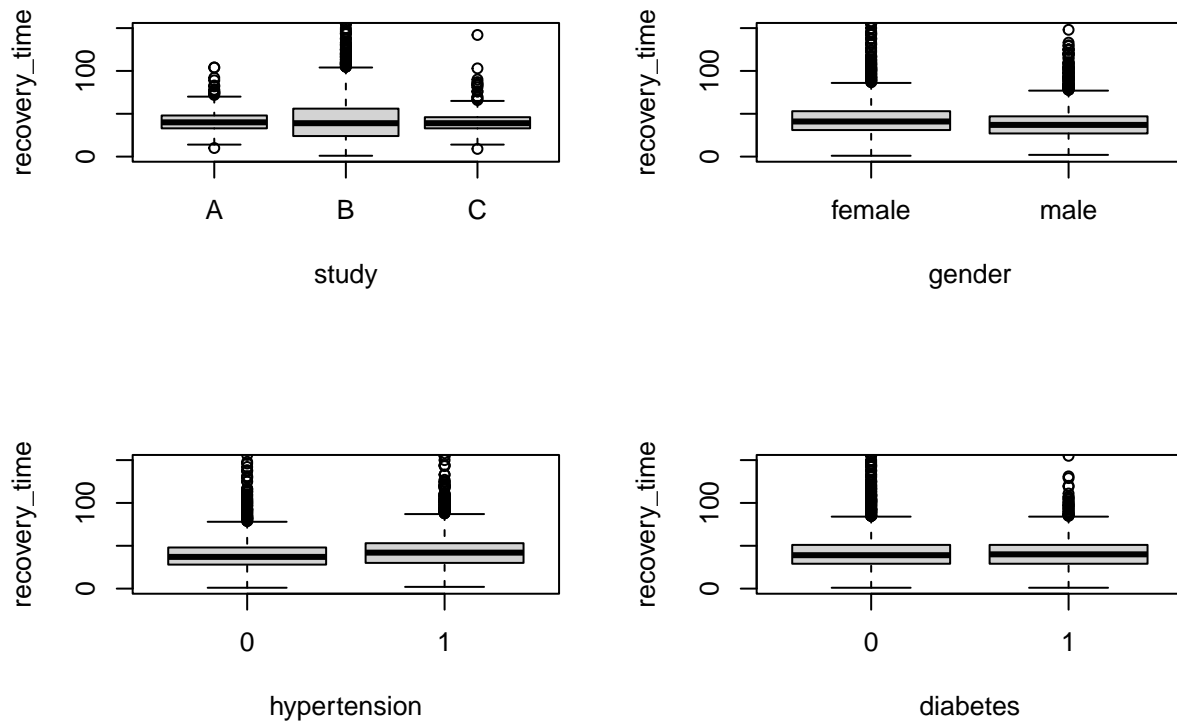
Figure 1. the relationship between predictors and recovery time



```

par(mfrow=c(2,2))
boxplot(recovery_time~study, data=dat, xlab="study", ylim=c(0,150))
boxplot(recovery_time~gender, data=dat, xlab="gender", ylim=c(0,150))
boxplot(recovery_time~hypertension, data=dat, xlab="hypertension", ylim=c(0,150))
boxplot(recovery_time~diabetes, data=dat, xlab="diabetes", ylim=c(0,150))





```



```

boxplot(recovery_time~vaccine, data=dat, xlab="vaccine", ylim=c(0,150))
boxplot(recovery_time~severity, data=dat, xlab="severity", ylim=c(0,150))
boxplot(recovery_time~race, data=dat, xlab="race", ylim=c(0,150))
boxplot(recovery_time~smoking, data=dat, xlab="smoking", ylim=c(0,150))

```

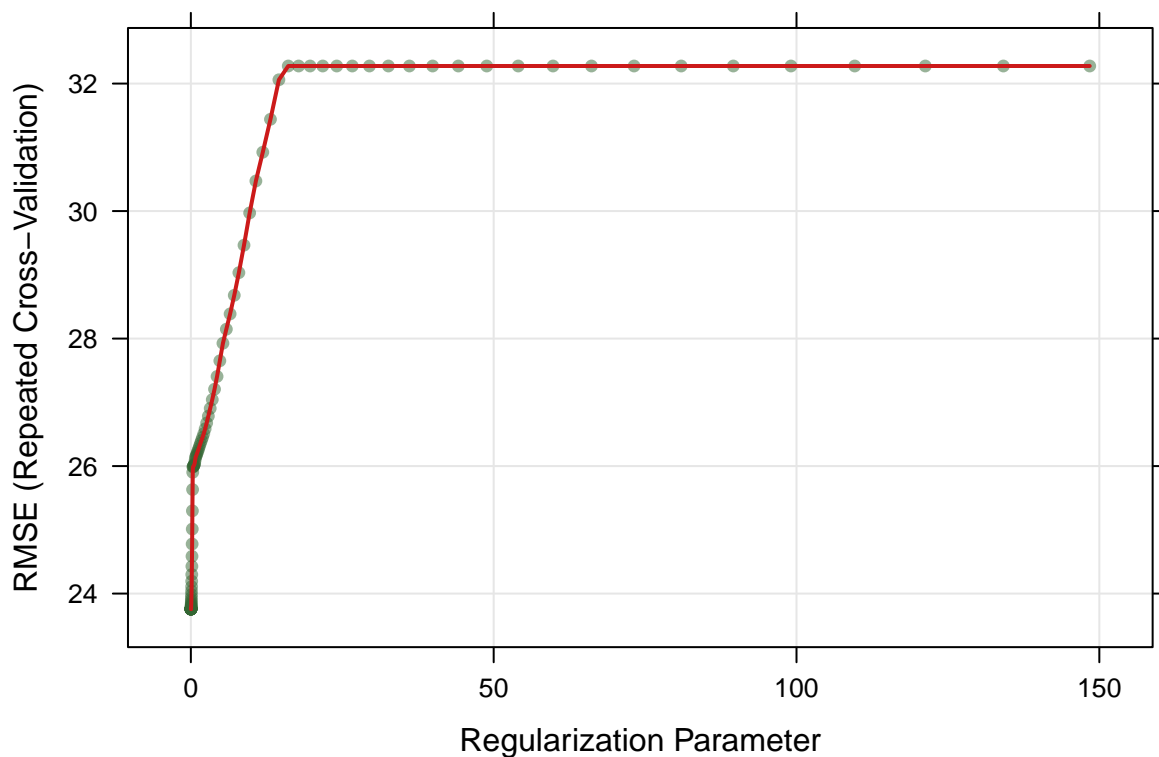


```
#linear model
set.seed(2023)
ctrl=trainControl(method = "repeatedcv", number =10, repeats = 5)
linear = train(recovery_time ~ age + gender + race + smoking + height +
               weight + bmi + hypertension + diabetes + sbp + ldl +
               vaccine + severity + study,
               data = train_dat,
               method = "lm",
               trControl = ctrl)
summary(linear$finalModel)
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -85.895 -14.897  -1.583   11.054  250.397
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.398e+03  1.372e+02 -24.774  < 2e-16 ***
## age          1.790e-01  1.272e-01   1.406   0.1597
## gendermale   -5.601e+00  1.008e+00  -5.556  3.02e-08 ***
## raceasian    -2.674e-01  2.354e+00  -0.114   0.9096
## raceblack    -2.943e+00  1.272e+00  -2.314   0.0207 *
## racehispanic -1.030e+00  1.734e+00  -0.594   0.5525
```

```
## smokingformer 5.233e+00 1.135e+00 4.611 4.19e-06 ***
## smokingcurrent 7.340e+00 1.696e+00 4.328 1.56e-05 ***
## height 1.973e+01 8.058e-01 24.479 < 2e-16 ***
## weight -2.134e+01 8.486e-01 -25.140 < 2e-16 ***
## bmi 6.424e+01 2.427e+00 26.468 < 2e-16 ***
## hypertension1 2.655e+00 1.682e+00 1.578 0.1146
## diabetes1 9.470e-01 1.373e+00 0.690 0.4904
## sbp 5.394e-02 1.096e-01 0.492 0.6226
## ldl -4.210e-02 2.686e-02 -1.567 0.1171
## vaccine1 -8.254e+00 1.028e+00 -8.025 1.46e-15 ***
## severity1 8.467e+00 1.630e+00 5.195 2.19e-07 ***
## studyB 6.723e+00 1.308e+00 5.139 2.95e-07 ***
## studyC 3.407e-01 1.595e+00 0.214 0.8309
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 26.97 on 2859 degrees of freedom
## Multiple R-squared: 0.3235, Adjusted R-squared: 0.3192
## F-statistic: 75.94 on 18 and 2859 DF, p-value: < 2.2e-16
#RMSE
test_pred1=predict(linear,newdata = test_dat)
rmse1=sqrt(mean((test_pred1-test_dat$recovery_time)**2))
rmse1
## [1] 23.74624
```


```
#lasso
set.seed(2023)
lasso=train(x1,y1,
            method = "glmnet",
            tuneGrid = expand.grid(alpha = 1,
                                   lambda = exp(seq(-5, 5, length = 100))),
            trControl = ctrl)
coef(lasso$finalModel, lasso$bestTune$lambda)
## 20 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -2.946231e+03
## age 1.216632e-01
## gendermale -3.415100e+00
## raceasian 1.035131e+00
## raceblack -2.068045e+00
## racehispanic -9.367240e-01
## smokingformer 3.432322e+00
## smokingcurrent 6.523183e+00
## height 1.698483e+01
## weight -1.833806e+01
## bmi 5.533357e+01
## hypertension1 1.243674e+00
## diabetes1 4.114282e-01
## sbp 4.419412e-02
## ldl -3.690833e-02
## vaccine1 -5.145263e+00
## severity1 5.402671e+00
## studyB 1.220249e+01
```

```
## studyC      .
## recovery_bins 2.949668e+01
lasso$bestTune$test_pred2=predict(lasso,newdata=x2)
pred_lasso=predict(lasso, newx = x2, s = lasso$lambda.min)
rmse_lasso= sqrt(mean((pred_lasso-y2)**2))
rmse_lasso
## [1] 34.30053
coef=coef(lasso, s = lasso$lambda.min)
n.pred=sum(coef[-1] != 0)
n.pred
## [1] 0
plot(lasso)
```



```
#elastic net
set.seed(2023)
elastic_net=train(x1, y1,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                         lambda = exp(seq(10, -10, length = 50))),
                  trControl = ctrl)
elastic_net$bestTune
##      alpha      lambda
## 411 0.4 0.002689588
test_pred_elastic=predict(elastic_net, newdata = x2)
rmse_elastic=sqrt(mean((test_pred_elastic - test_dat$recovery_time)**2))
```


```
rmse_elastic
## [1] 19.68248
plot(elastic_net)
```



```
#pls
set.seed(2023)
pls=train(x1, y1,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:15), # CHECK THIS
  trControl = ctrl,
  preProcess = c("center", "scale"))

summary(pls$finalModel)
## Data:      X dimension: 2878 19
## Y dimension: 2878 1
## Fit method: oscorespls
## Number of components considered: 15
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps
## X      8.425  15.83  24.05  30.50  34.82  41.85  46.44
## .outcome 31.437 34.49  34.93  35.15  35.36  35.67  38.49
##      8 comps 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X      48.98  52.73  54.76  59.54  65.03  70.23  75.72
## .outcome 44.60  46.18  47.24  47.37  47.37  47.37  47.37
##      15 comps
```

```
## X          80.18
## .outcome   47.37
test_pred_pls=predict(pls, newdata = x2)
rmse_pls=sqrt(mean((test_pred_pls - test_dat$recovery_time)**2))
rmse_pls
## [1] 19.75362
plot(pls)
```




```
#mars
set.seed(2023)
mars_grid = expand.grid(degree = 1:3,
                        nprune = 2:17)
mars = train(x1, y1,
             method = "earth",
             tuneGrid = mars_grid,
             trControl = ctrl)
kable(mars$bestTune, "simple")
```

	nprune	degree
48	17	3


```

coef(mars$finalModel)
##              (Intercept)              h(bmi-31.4)
##              21.054679              9.041933
##              recovery_bins              studyB * recovery_bins
##              20.720795              7.516211
##      h(bmi-31.4) * severity1 * studyB      h(age-63) * h(bmi-31.4) * studyB
##              20.672796              49.718869
##      h(bmi-26) * studyB * recovery_bins      h(26-bmi) * studyB * recovery_bins
##              3.331891              7.323239
##      h(bmi-32.2) * studyB * recovery_bins      h(bmi-31.4) * h(sbp-136) * studyB
##              50.562930              -3.451793
##      h(bmi-31.4) * h(136-sbp) * studyB      h(bmi-31.4) * h(ldl-119) * studyB
##              -1.480191              1.013023
##      vaccine1 * studyB * recovery_bins      smokingcurrent * h(bmi-31.4) * studyB
##              -8.417471              18.549723
##      h(age-64) * h(bmi-31.4) * studyB      h(age-61) * h(bmi-31.4) * studyB
##              -28.483749              -15.197017
##      gendermale * h(bmi-31.4) * studyB
##              -9.865810
test_pred_mars=predict(mars, newdata = x2)
rmse_mars=sqrt(mean((test_pred_mars - test_dat$recovery_time)**2))
rmse_mars
## [1] 15.78523
summary(mars)
## Call: earth(x=matrix[2878,19], y=c(15,56,42,62,4...), keepxy=TRUE, degree=3,
##              nprune=17)
##
##
##              coefficients
## (Intercept)              21.054679
## recovery_bins              20.720795
## h(bmi-31.4)              9.041933
## studyB * recovery_bins              7.516211
## vaccine1 * studyB * recovery_bins              -8.417471
## gendermale * h(bmi-31.4) * studyB              -9.865810
## smokingcurrent * h(bmi-31.4) * studyB              18.549723
## h(bmi-31.4) * severity1 * studyB              20.672796
## h(26-bmi) * studyB * recovery_bins              7.323239
## h(bmi-26) * studyB * recovery_bins              3.331891
## h(bmi-32.2) * studyB * recovery_bins              50.562930
## h(age-61) * h(bmi-31.4) * studyB              -15.197017
## h(age-63) * h(bmi-31.4) * studyB              49.718869
## h(age-64) * h(bmi-31.4) * studyB              -28.483749
## h(bmi-31.4) * h(sbp-136) * studyB              -3.451793
## h(bmi-31.4) * h(136-sbp) * studyB              -1.480191
## h(bmi-31.4) * h(ldl-119) * studyB              1.013023
##
## Selected 17 of 26 terms, and 10 of 19 predictors (nprune=17)
## Termination condition: Reached nk 39
## Importance: bmi, studyB, recovery_bins, age, severity1, sbp, ...
## Number of terms at each degree of interaction: 1 2 1 13
## GCV 244.2741    RSS 683133    GRSq 0.7714824    RSq 0.7777926
plot(mars)

```




```
#gam
gam = train(x1, y1,
            method = "gam",
            trControl = ctrl,
            control = gam.control(maxit = 200))
summary(gam$finalModel)
##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ gendermale + raceblack + racehispanic + smokingformer +
##   smokingcurrent + hypertension1 + diabetes1 + vaccine1 + severity1 +
##   studyB + studyC + recovery_bins + s(age) + s(sbp) + s(ldl) +
##   s(bmi) + s(height) + s(weight)
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   18.3728    1.4255  12.889  < 2e-16 ***
## gendermale    -2.9840    0.7735  -3.858  0.000117 ***
## raceblack     -1.2701    0.9647  -1.317  0.188071
## racehispanic  -0.6079    1.3147  -0.462  0.643825
## smokingformer  3.4597    0.8670   3.991  6.76e-05 ***
## smokingcurrent 7.1847    1.2936   5.554  3.05e-08 ***
## hypertension1  2.2592    0.7697   2.935  0.003361 **
## diabetes1      0.7835    1.0468   0.748  0.454233
```

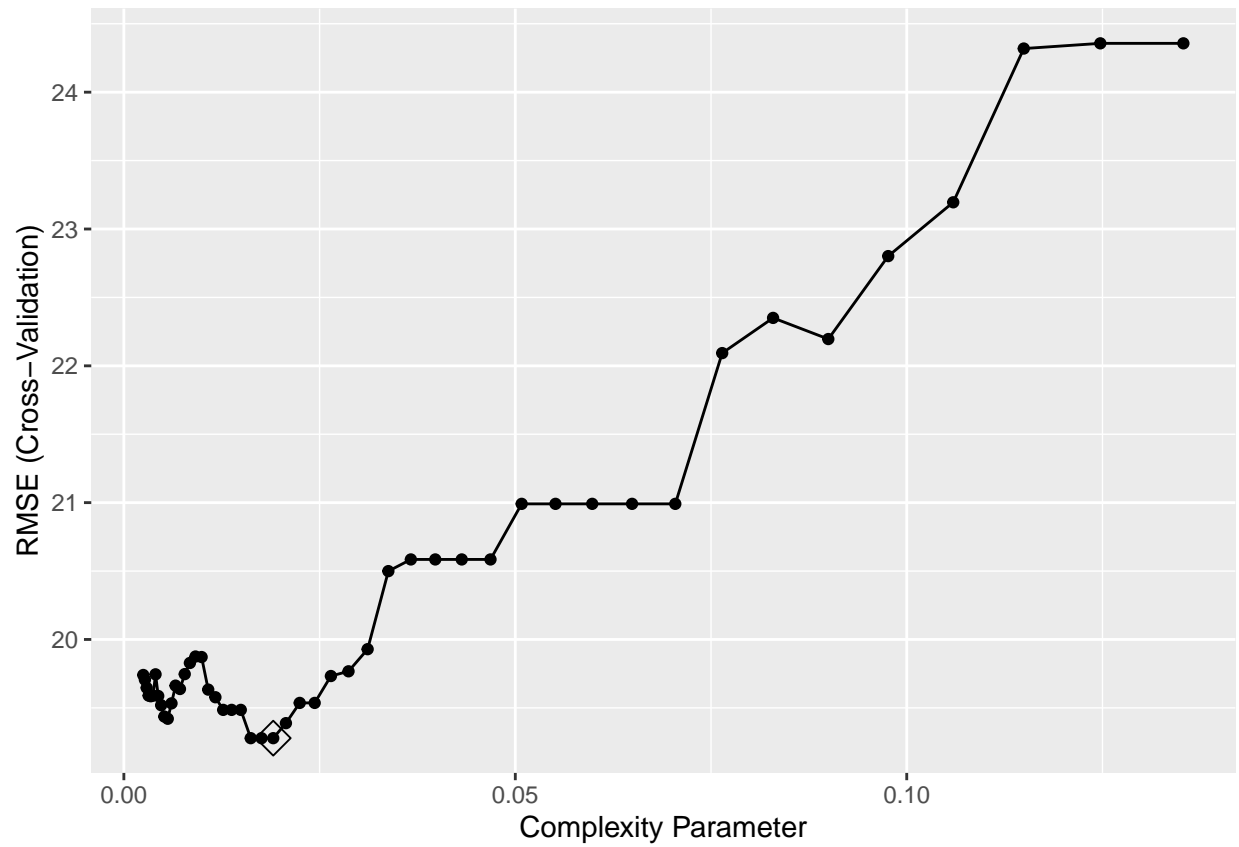
```

## vaccine1      -5.2527      0.7901  -6.648 3.54e-11 ***
## severity1     5.3175      1.2455   4.269 2.02e-05 ***
## studyB        12.3393      1.0120  12.193 < 2e-16 ***
## studyC         0.4666      1.2162   0.384 0.701252
## recovery_bins 28.1495      0.8977  31.358 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F  p-value
## s(age)       7.046e-08     9  0.000  0.692
## s(sbp)       4.882e-08     9  0.000  0.936
## s(ldl)       6.690e-08     9  0.000  0.393
## s(bmi)       6.939e+00     9 158.198 < 2e-16 ***
## s(height)    6.288e+00     9   4.214 1.04e-06 ***
## s(weight)    7.616e+00     9   6.804 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.606   Deviance explained =  61%
## GCV = 426.21  Scale est. = 421.2      n = 2878
gam$df.residual
## NULL
test_pred_gam=predict(gam, newdata = x2)
rmse_gam=sqrt(mean((test_pred_gam-test_dat$recovery_time)**2))
rmse_gam
## [1] 16.24292
plot(gam)

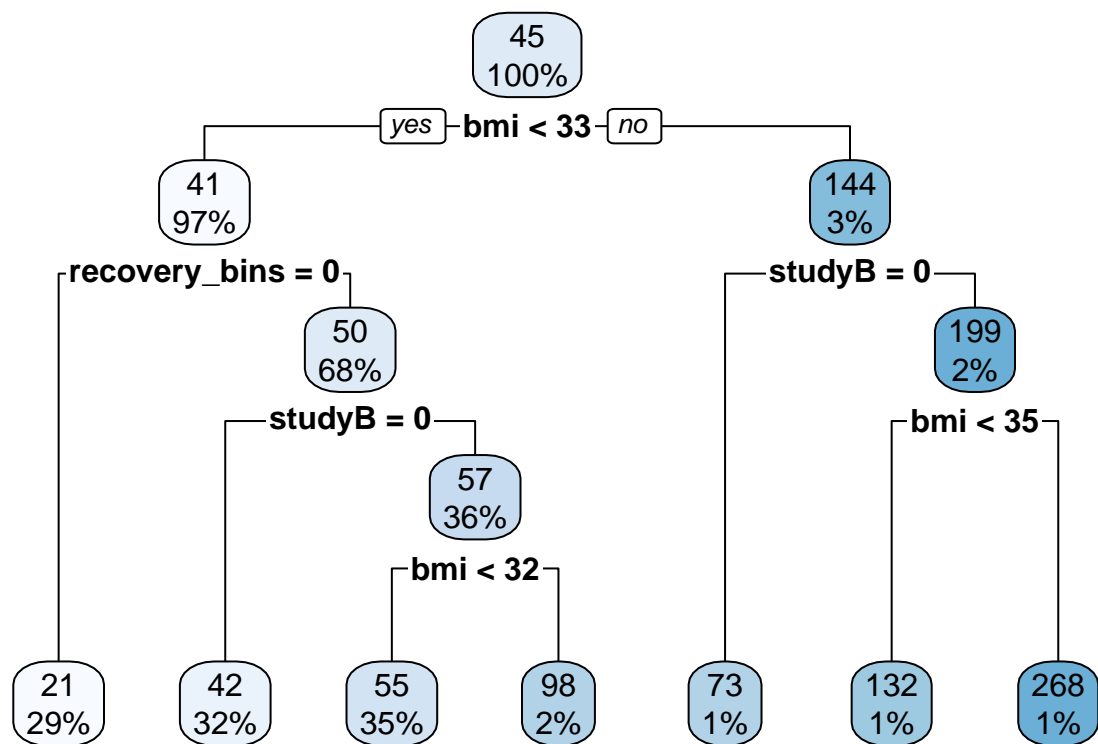
```



```
#tree
ctrl1=trainControl(method = "cv")
set.seed(2023)
rpart_fit <- train(recovery_time ~., data = dat[train_index,],
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-6, -2, length = 50))),
                  trControl = ctrl1)
rpart_fit$bestTune
##           cp
## 26 0.01907868
ggplot(rpart_fit, highlight = TRUE)
```



```
rpart.plot::rpart.plot(rpart_fit$finalModel)
```



The cp value is 0.0190787